

# Plant disease detection with CNNs: effects of real-time data augmentation and hyperparameters

Vision and Cognitive Services

---

Francesco Fontana   Lorenzo Mancini

03/09/2021

University of Padua

# Table of contents

---

1. Introduction
2. Dataset
3. Methods
4. Experiments
5. Conclusions

# Introduction

---

Introduction

Dataset

Methods

Experiments

Conclusions

# Goal

The **goal** of this project is to exploit computer vision resources in order to perform a detection for plant diseases and pests



**Figure 1:** Example of diseased plants: (from top left) potato, apple, strawberry, tomato

Plant automatic identification is a major land mark in agriculture:

- Disease identification positively impacts crop **yield** and **quality**
- It requires continuous monitoring by experts
- Historically, disease identification has been supported by agricultural extension organizations or other institutions, such as local plant clinics

Our project is organized as follows:

- Comparison between networks with same hyperparameters;
- Performances with data augmentation;
- Effects of hyperparameters;
- Main results and conclusions

# Dataset

---

Introduction

**Dataset**

Methods

Experiments

Conclusions

In order to achieve our goal, we make use of a public dataset available online **PlantVillage dataset** [2] 38 different classes and 54303 images (and more images are being collected)

The dataset is divided into 3 directory:

- **color** : contains original images in RGB
- **grayscale** : contains grayscaled version of original images
- **segmented** : RGB images with just the leaf segmented and color corrected.

For our goal we use just the **color** directory.



# Dataset example



**Figure 2:** *Example of colored (left), grayscaled (middle) and segmented samples (right).*

# Methods

---

Introduction

Dataset


Methods

Experiments

Conclusions

## Deep Learning methods:


- very efficient (high accuracy) and promising in the extracting features task [1];
- a lot of pre-trained models are available on the web and simple to load and customize;
- requires a huge amount of data for the training → computational heavy.



# Transfer Learning

Due to our hardware limitations, we implement **Transfer Learning** method and in particular we make use of **Feature Extraction** procedure:

- import pre-trained model (ex. VGG16/19, InceptionV3, etc.) with relative weights trained on the famous dataset “ImageNet”, without the top layers.
- freeze all parameters already present in the imported model so that they won't be trained again
- add top layers of the network adapted for our particular task (only these will be trained)



**Figure 3:** Example of *feature extraction* with VGG16's architecture

## Used models

In our project we've compared the results obtained with the following pre-trained models, already implemented on the *Tensorflow* library:

Model	n° of layers
VGG-16	16
VGG-19	19
InceptionV3	38
ResNet101V2	101

All these models can be tuned according to some parameters, such as: the **input size**, **epochs**, **the learning rate**, **the batch size**, the **dropout** and others.

# Hyper-Parameters


For all the previous models we firstly consider the following hyper-parameters as the “default” ones:

Epochs	30
Learning rate	0.001
Batch size	32
Dropout	0.15
Train – Test sizes	0.8 – 0.2

# Input size $i$

- First, we feed the models with different input sizes:


64x64



128x128




256x256



- See how much *input size* influences the overall performance and results of the models (ex. plot below with VGG16)

## Input size ii



**Figure 4:** Loss and accuracy plots of pre-trained VGG16 model with different image sizes in input

- **Note:** From now on all plots refers to 256x256 input size, since it provides the best results



## Results (“default” parameters)

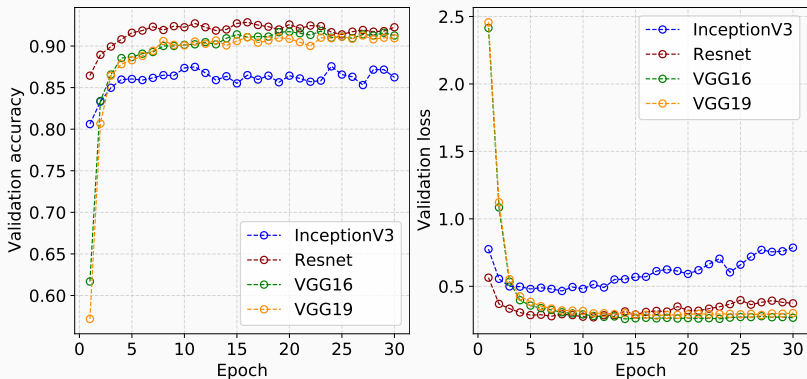
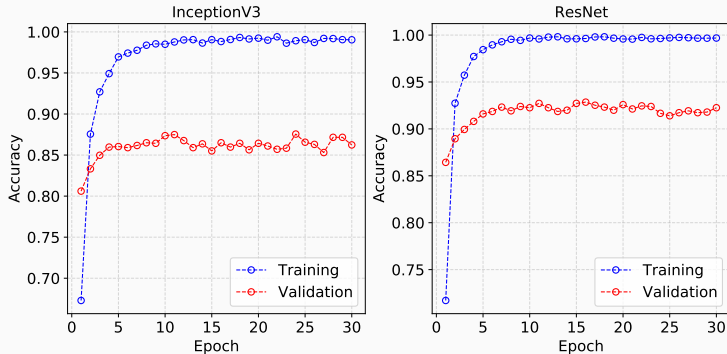


Figure 5: Results obtained with the “default” training.

# Overfitting (“default” parameters)

This standard approach leads in many cases to overfitting.



**Figure 6:** Accuracy plot of InceptionV3 and ResNet101V2 with “default” parameters. Here can be seen an evident example of overfitting

In order to solve the problem of overfitting, we try two different solutions:

- Real-time **data augmentation**;
- Tuning the hyperparameters.

# Experiments

---

Introduction

Dataset

Methods

**Experiments**

Conclusions


**Real-time data augmentation** allows to train the networks with more data without storing it in the memory:

- very efficient and easy to implement;
- reduces overfitting and improves performance;
- it represent a very good compromise when the input dataset is not so large.

It is implemented in the *Tensorflow* library:

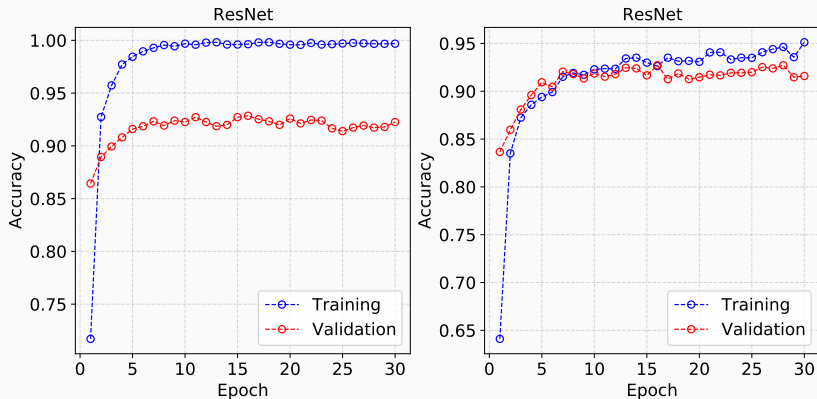
```
aug = ImageDataGenerator(  
    rotation_range=25,  
    width_shift_range=0.2,  
    height_shift_range=0.2,  
    shear_range=0.2,  
    zoom_range=0.3,  
    horizontal_flip=True,  
    fill_mode="nearest")
```

## Data Augmentation iii




**Figure 7:** *Examples of random transformations applied on-the-fly to images in order to perform data augmentation.*

# Data Augmentation iv



**Figure 8:** *Plots of loss and accuracy obtained during the training of the network for each model using the parameter*





**Figure 9:** Average difference between training and validation accuracy considering the last 10 epochs for all the networks.

## Batch size

**Batch size** is the number of samples that feed the network at each step: for example, a batch size = 32 means that the networks is trained taking 32 samples at time.

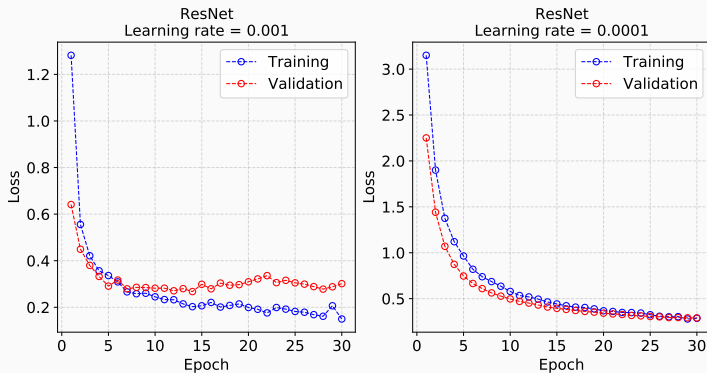
Model Name	Batch size	Accuracy
InceptionV3	32	87.6%
	120	87.5%
ResNet101V2	32	93.2%
	64	92.7%

**Table 1:** Comparison of the validation accuracy between different batch sizes: for InceptionV3 we use 32 and 120, whereas for ResNet101V2 we use 32 and 64. Performances are very similar.

**Learning rate (LR)** plays a major role in the loss optimization since it controls both how quickly the model is adapted to the problem and also the stability of the solutions

- **Low LR** (  $\leq 10^{-4}$  )
  - **Pro:** a more stable behaviour, less overfitting.
  - **Cons:** longer training time ( *more epochs required* )
- **High LR** (  $\sim 10^{-2}/10^{-3}$  )
  - **Pro:** Less computational time, good performance
  - **Cons:** more unstable solutions

## Learning Rate ii



**Figure 10:** *Plots of training and validation loss using 2 different learning rates. With the lower one we can see a more stable behaviour and less overfitting. On the other hand the higher learning rate provides a very good performance taking into account the fact that the computational time is lower.*

Model Name	Learning Rate	Accuracy
VGG16	1e-3	91.8%
	1e-4	89.5%
VGG19	1e-3	91.3%
	1e-4	89.6%
InceptionV3	1e-3	87.6%
	1e-4	87.4%
ResNet101V2	1e-3	93.2%
	1e-4	92.5%

**Table 2:** Comparison of the validation accuracy between different learning rate. (  $LR = 10^{-3} \rightarrow 30$  epochs;  $LR = 10^{-4} \rightarrow 50$  epochs)

# Dropout

To avoid the network to be a lot sensitive to each neurons' weight, we make use of **Dropout regularization** [3].

Model Name	Dropout	Accuracy
VGG16	0.15	90.9%
	0.5	91.1%
VGG19	0.15	90.5%
	0.5	90.7%
InceptionV3	0.15	87.7%
	0.5	87.1%
ResNet101V2	0.15	93.2%
	0.30	93.4%
	0.5	88.6%

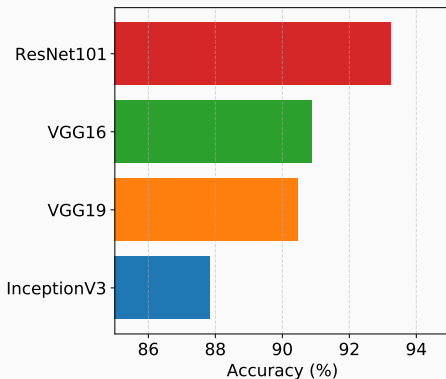
**Table 3:** Comparison of performance for all network with different dropouts

## Different train / test sizes

In order to further reduce the overfitting we've splitted the train and test sizes into a different way.

Model Name	Train – Test	Accuracy
VGG16	0.8 – 0.2	90.9%
	0.6 – 0.4	91.1%
VGG19	0.8 – 0.2	90.5%
	0.6 – 0.4	90.6%
InceptionV3	0.8 – 0.2	87.6%
	0.6 – 0.4	88.1%
ResNet101V2	0.8 – 0.2	93.2%
	0.6 – 0.4	92.9%

**Table 4:** Validation accuracy for InceptionV3 and ResNet101V2 splitting the entire dataset in different partition between training and test sets.



**Figure 11:** *Best performances obtained with each network. The highest accuracy (93.4%) is achieved by the ResNet model with data augmentation and dropout = 0.3.*



## Results ii

True label = **StrawberryLeafscorch**  
Predicted label = **StrawberryLeafscorch**  
Correctly classified = **True**  
Confidence = **1.0**



True label = **TomatoLeafMold**  
Predicted label = **TomatoLeafMold**  
Correctly classified = **True**  
Confidence = **0.917**



**Figure 12:** *Example of two correctly classified samples using ResNet101V2 model*

## Results iii

True label = **TomatoEarlyblight**  
Predicted label = **TomatoSpidermitesTwo**  
– **spottedspidermite**  
Correctly classified = **False**  
Confidence = **0.88**



True label = **TomatoSeptorialeafspot**  
Predicted label = **TomatoSpidermitesTwo**  
– **spottedspidermite**  
Correctly classified = **False**  
Confidence = **0.438**



**Figure 13:** *Example of two not correctly classified samples using ResNet101V2 model*

# Conclusions

---

Introduction

Dataset

Methods

Experiments

Conclusions

# Summary

---

In our work we make use of the public dataset **PlantVillage** in order to find a model able to predict diseases and pests. In particular we:

- exploit some pre-trained networks and study their performances and limitations;
- apply real-time **data augmentation** with the aim of reducing overfitting and improving performances;
- tune hyperparameters and study their impact on the networks.

At the end, we end up with our best performance equal to 93.4% accuracy provided by the ResNet101V2 network using data-augmentation and a different value of dropout.

# Hardware limitations

For our work we used the following hardware resources (referring to our local computer):

- **CPU:** AMD Ryzen 9 , 12 core, 24 threads (used multiprocessing of Tensorflow)
- **RAM memory:** 32 GB 3200 MHz




Figure 14: Training time for each network

## References i



S. M. Hassan, A. K. Maji, M. Jasiński, Z. Leonowicz, and E. Jasińska.  
**Identification of plant-leaf diseases using cnn and transfer-learning approach.**  
*Electronics*, 10(12), 2021.



D. P. Hughes and M. Salath'e .  
**An open access repository of images on plant health to enable the development of mobile disease diagnostics through machine learning and crowdsourcing.**  
*CoRR*, abs/1511.08060, 2015.



N. Srivastava, G. Hinton, A. Krizhevsky, I. Sutskever, and R. Salakhutdinov.  
**Dropout: A simple way to prevent neural networks from overfitting.**  
*Journal of Machine Learning Research*, 15(56):1929–1958, 2014.

**Thanks for your attention!**